

Experiment No. 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Steps:**1. Create an EC2 Linux Instances on AWS.**

[EC2](#) > [Instances](#) > Launch an instance

Launch an instance [Info](#)

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags [Info](#)

Name

[Add additional tags](#)

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Recents

Quick Start

Amazon Linux

aws

macOS

Ubuntu

ubuntu®


Windows

Microsoft

Red Hat

Red Hat

SUSE Li



[Browse more AMIs](#)

Including AMIs from

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

[Create new key pair](#)

▼ Network settings [Info](#)

[Edit](#)

Network [Info](#)

vpc-0404d393731afabc3

Subnet [Info](#)

No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)

Enable

Additional charges apply when outside of [free tier allowance](#)

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group

☐ Select existing security group

We'll create a new security group called 'launch-wizard-4' with the following rules:


☒ Allow SSH traffic from
Helps you connect to your instance

Anywhere

0.0.0.0/0

☒ Allow HTTPS traffic from the internet
To set up an endpoint, for example when creating a web server

☒ Allow HTTP traffic from the internet
To set up an endpoint, for example when creating a web server

 Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.



Instances (1) [Info](#)


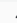



Last updated
less than a minute ago

[Connect](#)[Instance state ▼](#)[Actions ▼](#)[Launch instances ▼](#)

Find Instance by attribute or tag (case-sensitive)

Running ▼

< 1 > 

<input type="checkbox"/>	Name 	Instance ID	Instance state 	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4
<input type="checkbox"/>	abhinav	i-0a57da166d1f22307	 Running 	t2.micro	 2/2 checks passed View alarms +		us-east-1c	ec2-54-165-196-241.co...	54.165.196...

2. Then click on Id of that instance then click on connect

EC2 > Instances > i-0a57da166d1f22307 > Connect to instance

Connect to instance Info


Connect to your instance i-0a57da166d1f22307 (abhinav) using any of these options

EC2 Instance Connect

Session Manager


SSH client

EC2 serial console

**Port 22 (SSH) is open to all IPv4 addresses**

Port 22 (SSH) is currently open to all IPv4 addresses, indicated by **0.0.0.0/0** in the inbound rule in [your security group](#). For increased security, consider restricting access to only the EC2 Instance Connect service IP addresses for your Region: 18.206.107.24/29. [Learn more](#).

Instance ID


 i-0a57da166d1f22307 (abhinav)

Connection Type

☒ **Connect using EC2 Instance Connect**
Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address.



☐ **Connect using EC2 Instance Connect Endpoint**
Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.


Public IPv4 address

 54.165.196.241

Username

Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

 ec2-user 

**Note:** In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel **Connect**

3. SSH into the created machine instance

- Give permissions to the current user to the downloaded pem file using - `chmod 400 <security_filename.pem>`

```
abhin@Abhinavz-Acer MINGW64 ~ (master)
$ cd Downloads/

abhin@Abhinavz-Acer MINGW64 ~/Downloads (master)
$ chmod 400 "serverkey-01.pem"

abhin@Abhinavz-Acer MINGW64 ~/Downloads (master)
$ |
```

```
ssh -i (keyname).pem (username)@(public ipv4 dns address)
```

where keyname is name of the key you created. (server-01.pem). Other details can be found on the Instance dashboard.

[illegible]

4. Installation Of Docker007

➤ **sudo yum install docker -y**

```
[ec2-user@ip-172-31-90-103 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:13:41 ago on Sat Sep 14 03:42:27 2024.
Dependencies resolved.
=====
Package                Arch      Version                               Repository    Size
=====
Installing:
docker                 x86_64    25.0.6-1.amzn2023.0.2               amazonlinux   44 M
Installing dependencies:
containerd             x86_64    1.7.20-1.amzn2023.0.1               amazonlinux   35 M
iptables-libse         x86_64    1.8.8-3.amzn2023.0.2               amazonlinux   401 k
iptables-nft           x86_64    1.8.8-3.amzn2023.0.2               amazonlinux   183 k
libcgroup              x86_64    3.0-1.amzn2023.0.1                  amazonlinux   75 k
libnetfilter_conntrack x86_64    1.0.8-2.amzn2023.0.2               amazonlinux   58 k
libnfnfnetlink          x86_64    1.0.1-19.amzn2023.0.2              amazonlinux   30 k
libnftnl               x86_64    1.2.2-2.amzn2023.0.2               amazonlinux   84 k
pigz                   x86_64    2.5-1.amzn2023.0.3                  amazonlinux   83 k
runc                   x86_64    1.1.13-1.amzn2023.0.1              amazonlinux   3.2 M
Transaction Summary
-----
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libse-1.8.8-3.amzn2023.0.2.x86_64.rpm 2.2 MB/s | 401 kB    00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_64.rpm 2.5 MB/s | 183 kB    00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm 1.3 MB/s | 75 kB    00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64.rpm 1.3 MB/s | 58 kB    00:00
(5/10): libnfnfnetlink-1.0.1-19.amzn2023.0.2.x86_64.rpm 938 kB/s | 30 kB    00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rpm 1.6 MB/s | 84 kB    00:00
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm 1.7 MB/s | 83 kB    00:00
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm 21 MB/s | 3.2 MB    00:00
(9/10): containerd-1.7.20-1.amzn2023.0.1.x86_64.rpm 31 MB/s | 35 MB    00:01
(10/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rpm 28 MB/s | 44 MB    00:01
-----
Total                                                    52 MB/s | 84 MB    00:01
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing :
Installing : runc-1.1.13-1.amzn2023.0.1.x86_64
Installing : containerd-1.7.20-1.amzn2023.0.1.x86_64
Running scriptlet: containerd-1.7.20-1.amzn2023.0.1.x86_64
Installing : pigz-2.5-1.amzn2023.0.3.x86_64
Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
Installing : libnfnfnetlink-1.0.1-19.amzn2023.0.2.x86_64
Installing : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
Installing : iptables-libse-1.8.8-3.amzn2023.0.2.x86_64
Installing : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
```

```

Installed:
  containerd-1.7.20-1.amzn2023.0.1.x86_64      docker-25.0.6-1.amzn2023.0.2.x86_64
  libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64  libnfnetlink-1.0.1-19.amzn2023.0.2.x86_64

complete!
[ec2-user@ip-172-31-90-103 ~]$

```

- Configure cgroup in a daemon.json
(this can be done by creating the file and using **nano** text editor)

```

{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}

```

```

[ec2-user@ip-172-31-90-103 docker]$ sudo nano daemon.json
[ec2-user@ip-172-31-90-103 docker]$ |

```

```

ec2-user@ip-172-31-90-103:/etc/docker
GNU nano 5.8 daemon.json Modified
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}

```

- Enable and start docker and also load the daemon.json

```

sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker

```

```

[ec2-user@ip-172-31-90-103 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-90-103 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-90-103 docker]$ sudo systemctl restart docker
[ec2-user@ip-172-31-90-103 docker]$ |

```

- Check if docker is installed

```

[ec2-user@ip-172-31-90-103 docker]$ docker --version
Docker version 25.0.5, build 5dc9bcc
[ec2-user@ip-172-31-90-103 docker]$ |

```

5. Install Kubernetes

- SELinux needs to be disabled before configuring kubelet

sudo setenforce 0

sudo sed -i 's/^SELINUX=enforcing\$/SELINUX=permissive/' /etc/selinux/config

```
[ec2-user@ip-172-31-90-103 docker]$ sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-90-103 docker]$ |
```

- Add Kubernetes using the repo

(this is done by creating **kubernetes.repo** file in **/etc/yum.repos.d** and configuring it using **nano** editor)

[kubernetes]

name=Kubernetes

baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/

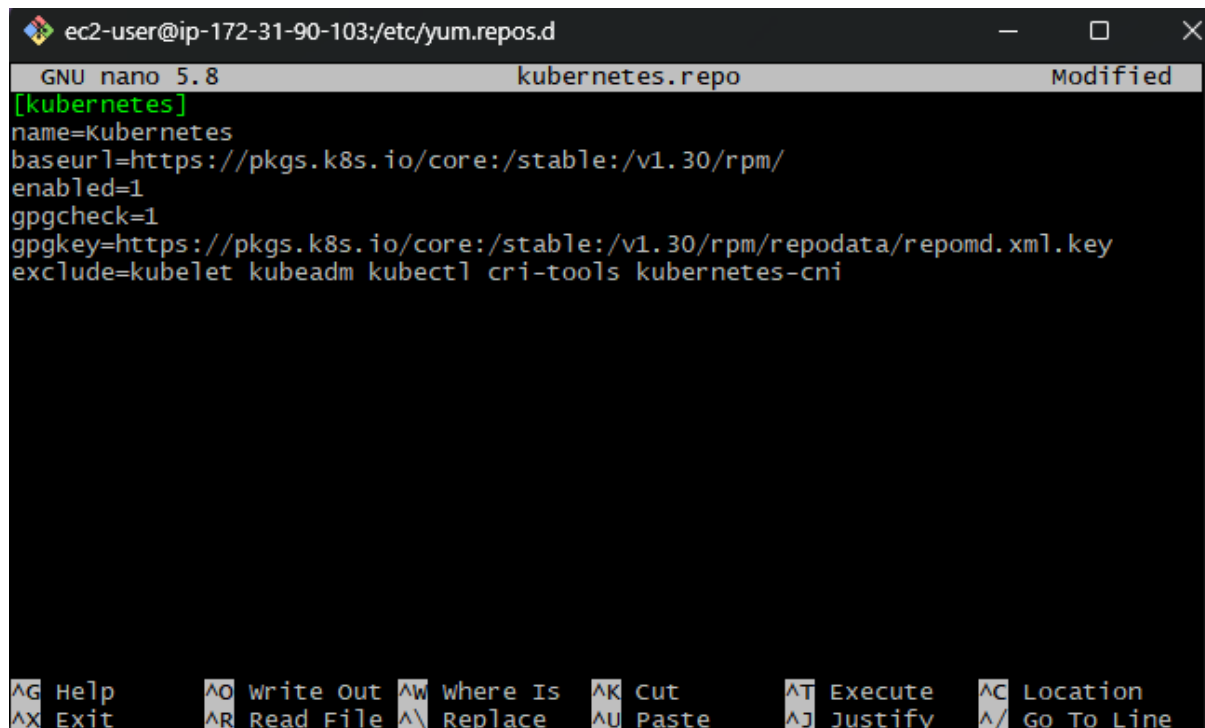
enabled=1

gpgcheck=1

gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key

exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni

```
[ec2-user@ip-172-31-90-103 docker]$ cd /etc/yum.repos.d/
[ec2-user@ip-172-31-90-103 yum.repos.d]$ ls
amazonlinux.repo  kernel-livepatch.repo
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo nano kubernetes.repo
[ec2-user@ip-172-31-90-103 yum.repos.d]$ ls
amazonlinux.repo  kernel-livepatch.repo  kubernetes.repo
[ec2-user@ip-172-31-90-103 yum.repos.d]$ |
```



The screenshot shows a terminal window with the nano text editor open at the file `/etc/yum.repos.d/kubernetes.repo`. The editor's title bar indicates it is GNU nano 5.8. The file content is as follows:

```
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

The bottom of the screen shows the nano editor's command shortcuts: `^G Help`, `^X Exit`, `^O Write Out`, `^R Read File`, `^W Where Is`, `^L Replace`, `^K Cut`, `^U Paste`, `^T Execute`, `^J Justify`, `^C Location`, and `^_ Go To Line`.

- Update packages list using **sudo yum update**

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo yum update
kubernetes                               125 kB/s | 17 kB      00:00
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```

- Install kubelet kubeadm kubectl

sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:00:42 ago on Sat Sep 14 04:08:20 2024.
Dependencies resolved.
=====
Package                        Arch      Version                               Repository    Size
=====
Installing:
kubeadm                        x86_64    1.30.5-150500.1.1                   kubernetes    10 M
kubectl                        x86_64    1.30.5-150500.1.1                   kubernetes    10 M
kubelet                        x86_64    1.30.5-150500.1.1                   kubernetes    17 M
Installing dependencies:
contrack-tools                 x86_64    1.4.6-2.amzn2023.0.2                 amazonlinux    208 k
cri-tools                      x86_64    1.30.1-150500.1.1                   kubernetes     8.6 M
kubernetes-cni                 x86_64    1.4.0-150500.1.1                   kubernetes     6.7 M
libnetfilter_cthelper          x86_64    1.0.0-21.amzn2023.0.2                 amazonlinux     24 k
libnetfilter_cttimeout         x86_64    1.0.0-19.amzn2023.0.2                 amazonlinux     24 k
libnetfilter_queue             x86_64    1.0.5-2.amzn2023.0.2                 amazonlinux     30 k
Transaction Summary
=====
Install 9 Packages

Total download size: 53 M
Installed size: 292 M
Downloading Packages:
(1/9): libnetfilter_cttimeout-1.0.0-19.amzn2023 448 kB/s | 24 kB    00:00
(2/9): libnetfilter_cthelper-1.0.0-21.amzn2023. 409 kB/s | 24 kB    00:00
(3/9): libnetfilter_queue-1.0.5-2.amzn2023.0.2. 1.5 MB/s | 30 kB    00:00
(4/9): contrack-tools-1.4.6-2.amzn2023.0.2.x86 1.8 MB/s | 208 kB    00:00
(5/9): cri-tools-1.30.1-150500.1.1.x86_64.rpm 28 MB/s | 8.6 MB    00:00
(6/9): kubectl-1.30.5-150500.1.1.x86_64.rpm 23 MB/s | 10 MB    00:00
(7/9): kubeadm-1.30.5-150500.1.1.x86_64.rpm 18 MB/s | 10 MB    00:00
(8/9): kubelet-1.30.5-150500.1.1.x86_64.rpm 37 MB/s | 17 MB    00:00
(9/9): kubernetes-cni-1.4.0-150500.1.1.x86_64.r 20 MB/s | 6.7 MB    00:00
-----
Total                                           56 MB/s | 53 MB    00:00
```

```
Installed:
contrack-tools-1.4.6-2.amzn2023.0.2.x86_64
cri-tools-1.30.1-150500.1.1.x86_64
kubeadm-1.30.5-150500.1.1.x86_64
kubectl-1.30.5-150500.1.1.x86_64
kubelet-1.30.5-150500.1.1.x86_64
kubernetes-cni-1.4.0-150500.1.1.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```

- After installing Kubernetes, we need to configure internet options to allow bridging.

sudo swapoff -a

echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf

sudo sysctl -p

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ . sudo swapoff -a
-bash: sudo: No such file or directory
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo swapoff -a
[ec2-user@ip-172-31-90-103 yum.repos.d]$ echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo sysctl -p
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-90-103 yum.repos.d]$ |
```


6. Initialize the Kubecluster

`sudo kubeadm init --podnetwork-cidr=10.244.0.0/16`

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all
I0914 04:12:17.448521 27990 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks
        [WARNING NumCPU]: the number of available CPUs 1 is less than the required 2
        [WARNING Mem]: the system RAM (949 MB) is less than the minimum 1700 MB
        [WARNING FileExisting-socat]: socat not found in system path
        [WARNING FileExisting-tc]: tc not found in system path
        [WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl enable kubelet.service'
[preflight] pulling images required for setting up a kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
W0914 04:12:17.711154 27990 checks.go:844] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.9" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
```

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.90.103:6443 --token 0zk8w3.xyegkydsy42vfscm \
--discovery-token-ca-cert-hash sha256:31c672892b19dcb869fc46362d189234128f5bfc302bd41ae8c6078c56173f00
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```

➤ Save the token

```
kubeadm join 172.31.90.103:6443 --token 0zk8w3.xyegkydsy42vfscm \
--discovery-token-ca-cert-hash sha256:31c672892b19dcb869fc46362d189234128f5bfc302bd41ae8c6078c56173f00
```

➤ Copy the mkdir and chown commands from the top and execute them

`mkdir -p $HOME/.kube`

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ mkdir -p $HOME/.kube
[ec2-user@ip-172-31-90-103 yum.repos.d]$ |
```

`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```

`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```


- Then, add a common networking plugin called flannel file as mentioned in the code.

kubectl apply -f <https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[ec2-user@ip-172-31-90-103 yum.repos.d]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[ec2-user@ip-172-31-90-103 yum.repos.d]$
```

7. Deploy nginx server

- Apply deployment using this following command:

kubectl apply -f <https://k8s.io/examples/pods/simple-pod.yaml>

```
[ec2-user@ip-172-31-90-103 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
```

- use **kubectl get nodes** to check whether the pod gets created or not

```
[ec2-user@ip-172-31-90-103 docker]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0          12s
```

- To convert state from pending to running use following command:

kubectl describe pod nginx (This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.)

```
[ec2-user@ip-172-31-90-103 docker]$ kubectl describe pod nginx
Name:      nginx
Namespace: default
Priority:   0
Service Account: default
Node:      <none>
Labels:    <none>
Annotations: <none>
Status:    Pending
IP:        <none>
IPs:       <none>
Containers:
  nginx:
    Image:      nginx:1.14.2
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-k4lj6 (ro)
  Type:      Projected (a volume that contains injected data from multiple sources)
  TokenExpirationSeconds: 3607
  ConfigMapName: kube-root-ca.crt
  ConfigMapOptional: <nil>
  DownwardAPI: true
QoS Class:   BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
             node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason             Age   From          Message
  ----     -
  Warning   FailedScheduling   7s    default-scheduler 0/1 nodes are available: 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }. Preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
```

- check pod status

```
[ec2-user@ip-172-31-90-103 ~]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   1 (6s ago)  90s
```

- mention the port you want to host

```
[ec2-user@ip-172-31-90-103 ~]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

8. Verify your deployment

- Open up a new terminal and ssh to your EC2 instance. Then, use this curl command to check if the Nginx server is running. **curl --head http://127.0.0.1:8080** If the response is 200 OK and you can see the Nginx server name, your deployment was successful. We have successfully deployed our Nginx server on our EC2 instance.

```
[ec2-user@172-31-90-103 ~]$ curl --head http://127.0.0.1:8080
HTTP/1.1 200 OK
Server: nginx/1.18.0 (Linux)
Date: Sat, 14 Sep 2024 12:31:53 GMT
Content-Type: text/html
Content-Length: 612
Connection: keep-alive
```

Conclusion:

An AWS EC2 Linux instance was set up, and Docker and Kubernetes were installed. Kubernetes was initialized successfully, and the required commands were executed. Flannel was installed as a networking plugin. Although there was an initial error with the Nginx deployment, it was eventually deployed successfully using the `simple-pod.yml` file and accessed via localhost on port 8080.